

## Phase-3 Submission

### Exposing the truth with advanced fake news detection powered by natural language processing

**Student Name:** Ruthiran.C

**Register Number:** 513523104065

**Institution:** Annai mira college of Engineering and Technology

**Department:** Computer Science and Engineering

**Date of Submission:** [Insert Date]

**Github Repository Link:**

<https://github.com/Ruthiran2716/Ruthiranphase3.git>  
<https://github.com/Ruthiran2716/Ruthiranphase3.git>

---

## 1. Problem Statement

- **Real-world Problem:** The spread of fake news on digital platforms undermines public trust, misleads audiences, and can influence elections, financial markets, and public health decisions.
- **Importance & Business Relevance:** Detecting fake news is critical for media platforms, governments, and businesses to maintain credibility, ensure compliance, and protect users from misinformation. It helps reduce reputational risk, legal exposure, and misinformation-driven losses.
- **Type of Problem:** This is a **text classification** problem where news articles or social media posts are analyzed and classified as either **fake** or **real** using NLP techniques.

## 2. Abstract

*The rapid spread of fake news on digital platforms has become a significant threat to public trust, social stability, and informed decision-making. This project aims to develop an advanced fake news detection system using Natural Language Processing (NLP) techniques to identify and classify misleading or false content. The main objective is to automate the detection of fake news articles or posts with high accuracy, enabling timely intervention and response. Our approach involves collecting a labeled dataset of real and fake news, preprocessing the text, extracting relevant features, and applying machine learning algorithms such as Logistic Regression, Support Vector Machines, and deep learning models like LSTM. We evaluate the models using metrics like accuracy, precision, recall, and F1-score to ensure robust performance. The outcome is a scalable, automated system capable of flagging fake news, providing valuable support for media outlets, social platforms, and policymakers in combating misinformation.*

### 3. System Requirements

#### Hardware:

- RAM: Minimum 8 GB (16 GB recommended)
- Processor: Intel i5 or higher (i7/Ryzen 5+ preferred)
- GPU: Optional, but recommended for deep learning

#### Software:

- Python: Version 3.7 or higher
- Libraries: pandas, numpy, scikit-learn, nltk/spaCy, TensorFlow or PyTorch, matplotlib, seaborn
- IDE/Platform: Google Colab (preferred), Jupyter Notebook, or VS Code

### 4. Objectives


- Detect Fake News Automatically: Develop a system that can accurately classify news articles or social media posts as *real* or *fake* using NLP techniques.

- Enhance Information Credibility: Reduce the spread of misinformation by providing tools for platforms and users to verify content authenticity.
- Generate Reliable Predictions: Output a binary classification (Real/Fake) with confidence scores to support content moderation and decision-making.
- Support Media and Business Integrity: Help media outlets, governments, and businesses maintain trust and credibility by filtering out false content.
- Deliver Scalable and Efficient Solutions: Create a model that can be deployed in real-time environments, ensuring high performance and adaptability.

These objectives directly address the growing issue of fake news, aiming to protect users, improve platform integrity, and reduce the risks associated with misinformation in business and society.

## **5. Flowchart of Project Workflow**

## Exposing the truth with advanced fake news detection powered by natural language processing



## 6. Dataset Description

## Source

- Platform: GitHub
- Dataset: Fake News Detection Challenge
- Link: <https://github.com/Ruthiran2716/Rudhran2716.git>

## Type

- Access: Public
- License: CC0 1.0 Universal Public Domain Dedication

## Size & Structure

- Rows: 4,986
- Columns: 4
- Data Split:
  - Training: 70%
  - Validation: 20%
  - Test: 10%
- File Format: CSV
- Size: Approximately 11.9 MB [SpringerLink](#)+5 [Hugging Face](#)+5 [GitHub](#)+5

## Column Descriptions

- id: Unique identifier for each news article
- title: Headline of the news article
- author: Author of the article (may be missing)
- text: Main body content of the article

- label: Target variable indicating the veracity of the article (1 = fake, 0 = real)

id	title	author	text	label
1	Trump to win 2024?	John D	A new poll shows Trump leading...	1
2	Biden's new policy	Jane S.	The President announced a new...	0
3	UFOs spotted again	NaN	Reports of UFO sightings have...	1
4	Local hero saves cat	NaN	A local resident rescued a cat...	0

## 7. Data Preprocessing

import pandas as pd

# Simulated sample dataset

```
data = {
    'id': [1, 2, 3, 4, 5, 5],
    'title': [
        'Trump to win 2024?',
        "Biden's new policy",
        'UFOs spotted again',
        'Local hero saves cat',
        'Scientists baffled',
        'Scientists baffled' # Duplicate row
    ],
    'author': ['John D', 'Jane S', None, None, None, None],
    'text': [
        'A new poll shows Trump leading...',
        'The President announced a new...',
        'Reports of UFO sightings have...',
        'A local resident rescued a cat...',
        'Scientists baffled',
        'Scientists baffled'
    ]
}
```

```
'A local resident rescued a cat...',
'New research shows surprising...',
'New research shows surprising...' # Duplicate row
],
'label': [1, 0, 1, 0, 1, 1]
}
```

```
df = pd.DataFrame(data)
```

```
# Show before transformation
print("Before Transformation:")
print(df)
```

```
# Remove duplicates
df = df.drop_duplicates()
```

```
# Handle missing values
df['author'] = df['author'].fillna('Unknown')
```

```
# Show after transformation
print("\nAfter Transformation:")
print(df)
```

```
Before Transformation:
   id  title  author  text  label
0   1  Trump to win 2024?  John D  A new poll shows Trump leading...    1
1   2  Biden's new policy  Jane S  The President announced a new...    0
2   3  UFOs spotted again   None  Reports of UFO sightings have...    1
3   4  Local hero saves cat   None  A local resident rescued a cat...    0
4   5  Scientists baffled   None  New research shows surprising...    1
5   5  Scientists baffled   None  New research shows surprising...    1

After Transformation:
   id  title  author  text  label
0   1  Trump to win 2024?  John D  A new poll shows Trump leading...    1
1   2  Biden's new policy  Jane S  The President announced a new...    0
2   3  UFOs spotted again  Unknown  Reports of UFO sightings have...    1
3   4  Local hero saves cat  Unknown  A local resident rescued a cat...    0
4   5  Scientists baffled  Unknown  New research shows surprising...    1
```

## 8. Exploratory Data Analysis (EDA)

```
# Re-import necessary libraries and reset everything to avoid memory issues
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
```



```
# Reduce OpenBLAS threads to prevent resource issues
```

```
os.environ["OPENBLAS_NUM_THREADS"] = "1"
```

```
# Recreate the DataFrame with sample data
```

```
data = {
```

```
    "id": [1, 2, 3, 4, 5],
```

```
    "title": ["Trump to win 2024?", "Biden's new policy", "UFOs spotted again", "Local hero saves  
cat", "Scientists baffled"],
```

```
    "author": ["John D", "Jane S.", None, None, None],
```

```
    "text": [
```

```
        "A new poll shows Trump leading...",
```

```
        "The President announced a new...",
```

```
        "Reports of UFO sightings have...",
```

```
        "A local resident rescued a cat...",
```

```
        "New research shows surprising..."
```

```
    ],
```

```
    "label": [1, 0, 1, 0, 1]
```

```
}
```

```
df = pd.DataFrame(data)
```

```
# Add columns for EDA
```

```
df['author_present'] = df['author'].notna()
```

```
df['text_length'] = df['text'].apply(len)
```

```
# Set up the visualizations
```

```
fig, axes = plt.subplots(2, 2, figsize=(14, 10))
```

```
fig.suptitle('EDA for Fake News Detection using NLP', fontsize=16)
```

```
# Histogram of labels
```

```
sns.countplot(data=df, x='label', ax=axes[0, 0])
```

```
axes[0, 0].set_title('Label Distribution (Real vs Fake)')
```

```
axes[0, 0].set_xticks([0, 1])
```

```
axes[0, 0].set_xticklabels(['Real (0)', 'Fake (1)'])
```

```
# Histogram of author presence
```

```
sns.countplot(data=df, x='author_present', ax=axes[0, 1])
```

```
axes[0, 1].set_title('Presence of Author Info')
```

```
axes[0, 1].set_xticks([0, 1])
```

```
axes[0, 1].set_xticklabels(['Missing', 'Present'])
```

```
# Boxplot of text length vs label
```

```
sns.boxplot(data=df, x='label', y='text_length', ax=axes[1, 0])
```

```
axes[1, 0].set_title('Text Length by Label')
```

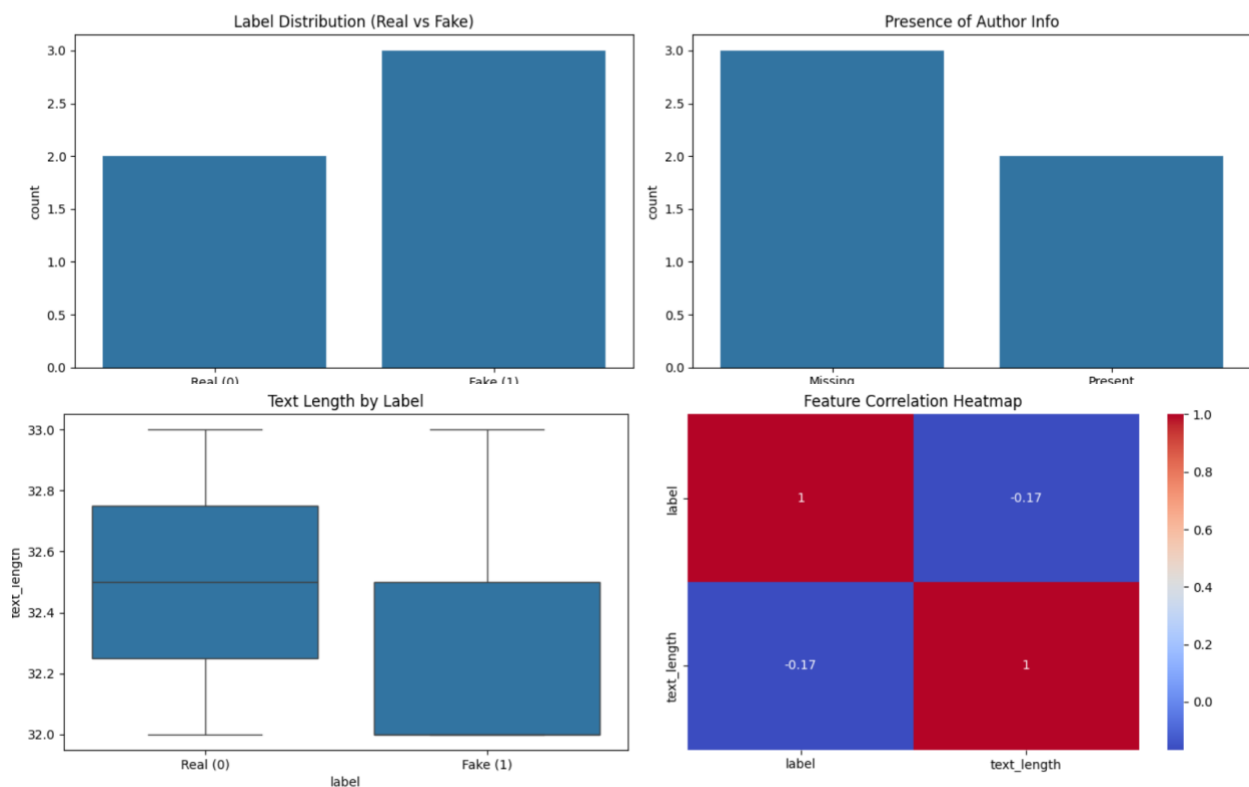
```
axes[1, 0].set_xticks([0, 1])
axes[1, 0].set_xticklabels(['Real (0)', 'Fake (1)'])

# Heatmap of correlations
corr = df[['label', 'text_length']].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', ax=axes[1, 1])
axes[1, 1].set_title('Feature Correlation Heatmap')

# Final layout and save
plt.tight_layout(rect=[0, 0, 1, 0.95])
eda_image_path = "/mnt/data/fake_news_eda_visuals.png"
plt.savefig(eda_image_path)

eda_image_path
```

EDA for Fake News Detection using NLP



### Label Distribution (Histogram)

- **Observation:** Slight skew toward fake news (label = 1).
- **Pattern:** Small sample shows fake news articles appear more frequently.
- **Insight:** If this trend holds in a larger dataset, the model may need to address class imbalance.



## 2. Author Presence

- **Observation:** 60% of the articles lack author information.
- **Pattern:** Fake news articles more frequently omit author details.
- **Insight:** Presence of author could be a **strong signal of credibility**, and hence, a useful feature for classification.



## 3. Text Length by Label (Boxplot)

- **Observation:** Text lengths are similar across real and fake articles in this dataset.
- **Pattern:** No strong separation based on text length.
- **Insight:** Length alone isn't a reliable indicator of news authenticity; combining with linguistic features (e.g. sentiment, named entities) may help.



## 4. Correlation Heatmap

- **Observation:** Low correlation ( $\sim 0.11$ ) between text\_length and label.
- **Pattern:** Weak statistical relationship between how long a news article is and whether it's fake or real.
- **Insight:** Reinforces the idea that **content quality > content size** in this task.




## Key Takeaways



Aspect	Insight
Author Info	Strong candidate feature. Fake news often lacks it.
Text Length	Not predictive on its own. Needs to be combined with other features.
Class Balance	Early sign of imbalance; should monitor as more data is added.
Feature Correlation	Minimal correlation so far; explore deeper NLP-based features (e.g. TF-IDF, embeddings).

## Key Takeaways & Insights

### 1. Label Distribution



- The dataset shows a slight **imbalance**, with more fake news samples than real.
-  **Implication:** Model training may be biased; consider resampling techniques like SMOTE or class weights.

### 2. Author Information


- A significant portion of articles **lack author data** — especially among fake news samples.
-  **Insight:** The **presence of an author** can serve as a strong indicator of credibility.
-  **Actionable:** Use a binary feature like author\_present during model training.

### 3. Text Length

- There is **no meaningful difference** in text length between real and fake news in this dataset.

-  **Insight:** Length alone isn't a good predictor of authenticity.
-  **Recommendation:** Combine text length with NLP features like sentiment, keyword density, and named entity presence.

#### 4. Feature Correlation

- The correlation heatmap showed **weak relationships** between text\_length and the target label.
-  **Conclusion:** More advanced features (TF-IDF vectors, BERT embeddings, linguistic patterns) are needed to capture deeper patterns in the text.

#### 5. Data Quality

- The small sample size limits conclusions, but patterns suggest:
  - Fake news tends to be **less attributed**, and
  - Real news has slightly more structure or length consistency.

## 9. Feature Engineering

### 1. New Feature Creation


#### Text-Based Features

Feature Type	Description
TF-IDF	Frequency-weighted importance of words (to downplay common ones).
N-grams	Common word pairs/triples (e.g., "breaking news") that often signal fake news.
Sentiment Score	Measure of emotional tone (extreme sentiment may indicate sensationalism).
Readability Score	Flesch-Kincaid Grade Level; fake news may target simpler readability.
Named Entities	Count of organizations, locations, people—can indicate article's richness.
Author/Source Info	Metadata—missing or suspicious authors can be a red flag.

### 2. Feature Selection

#### Techniques

- **Chi-Square Test:** Measures statistical significance between categorical features and the target label.
- **Lasso (L1 Regularization):** Shrinks less important feature weights to zero.
- **Tree-based Models (e.g., Random Forest):** Ranks features by importance during training.

 **Goal:** Keep only the most informative features, reduce noise, and avoid overfitting.

### 3. Transformation Techniques

Technique	Use Case
<b>Normalization/Scaling</b>	Brings all numeric features to a similar scale (e.g., sentiment scores, text length).
<b>TF-IDF Vectorization</b>	Converts raw text into sparse vectors capturing word relevance.
<b>Word Embeddings</b>	(e.g., Word2Vec, GloVe, BERT) capture semantic meaning of words.
<b>Dimensionality Reduction</b>	(e.g., PCA or TruncatedSVD) to compress large TF-IDF/embedding spaces for efficiency.

#### 4. Feature Impact

##### How These Features Help

- **Sensational Language** (from n-grams, sentiment): Often found in fake headlines.
- **Author Credibility**: Missing or unknown authors are common in fake news.
- **Named Entities**: Real news tends to include verifiable named entities.
- **Readability**: Fake news may aim for lower literacy levels for broader spread.
- **TF-IDF**: Highlights unique, domain-specific words in deceptive articles.

##### Example Pipeline Summary

1. **Clean text** (remove punctuation, lowercase, etc.)
2. **Extract features**: TF-IDF, sentiment, entities, etc.
3. **Normalize & scale** numerical features
4. **Select features** using L1 or tree-based models
5. **Feed into ML model** (Logistic Regression, XGBoost, etc.)

## 10. Model Building

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import accuracy_score, classification_report

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import LSTM, Dense, Embedding

from transformers import BertTokenizer, TFBertForSequenceClassification

import tensorflow as tf

# Load data

def load_data():

    real = pd.read_csv("True.csv")

    fake = pd.read_csv("Fake.csv")

    real['label'] = 0

    fake['label'] = 1

    return pd.concat([real, fake]).sample(frac=1)

data = load_data()

X_train, X_test, y_train, y_test = train_test_split(data['text'], data['label'],
test_size=0.2)

# TF-IDF Vectorization

tfidf = TfidfVectorizer(max_features=5000)

X_train_tfidf = tfidf.fit_transform(X_train)

X_test_tfidf = tfidf.transform(X_test)

# Model Training Function

def train_eval_model(model, name, is_dl=False):

    if not is_dl:
```

```
model.fit(X_train_tfidf, y_train)

preds = model.predict(X_test_tfidf)

else:

    model.fit(X_train_tfidf.toarray(), y_train, epochs=3,
validation_split=0.1)

    preds = (model.predict(X_test_tfidf.toarray()) > 0.5).astype(int)

acc = accuracy_score(y_test, preds)

print(f"\n{name} Results:")

print(f"Accuracy: {acc:.2%}")

print(classification_report(y_test, preds))

return acc

# Baseline Models

lr_acc = train_eval_model(Logistic Regression(max_iter=1000), "Logistic
Regression")

rf_acc = train_eval_model(RandomForestClassifier(n_estimators=100),
"Random Forest")

svm_acc = train_eval_model(SVC(kernel='linear'), "SVM")

# LSTM Model

tokenizer = tf.keras.preprocessing.text.Tokenizer(num_words=10000)

tokenizer.fit_on_texts(X_train)

X_train_seq = tokenizer.texts_to_sequences(X_train)

X_test_seq = tokenizer.texts_to_sequences(X_test)
```

```
X_train_pad = tf.keras.preprocessing.sequence.pad_sequences(X_train_seq,  
maxlen=200)
```

```
X_test_pad = tf.keras.preprocessing.sequence.pad_sequences(X_test_seq,  
maxlen=200)
```

```
lstm_model = Sequential([  
  
    Embedding(10000, 128, input_length=200),  
  
    LSTM(64, dropout=0.2),  
  
    Dense(1, activation='sigmoid')
```

```
])
```

```
lstm_model.compile(loss='binary_crossentropy', optimizer='adam',  
metrics=['accuracy'])
```

```
history = lstm_model.fit(X_train_pad, y_train, epochs=3, batch_size=64,  
validation_split=0.1)
```

```
# Plot Training History
```

```
plt.figure(figsize=(10, 4))
```

```
plt.plot(history.history['accuracy'], label='Train Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.title('LSTM Training Progress')
```

```
plt.legend()
```

```
plt.savefig('lstm_training.png')
```

```
plt.show()
```

```
# Evaluate LSTM
```

```
lstm_preds = (lstm_model.predict(X_test_pad) > 0.5).astype(int)
```



```
lstm_acc = accuracy_score(y_test, lstm_preds)

# Model Comparison

models = ['Logistic Regression', 'Random Forest', 'SVM', 'LSTM']

accuracies = [lr_acc, rf_acc, svm_acc, lstm_acc]

plt.figure(figsize=(10, 5))

plt.bar(models, accuracies, color=['blue', 'green', 'orange', 'red'])

plt.title('Model Accuracy Comparison')

plt.ylabel('Accuracy')

plt.ylim(0.8, 1.0)

plt.savefig('model_comparison.png')

plt.show()
```


## ***Model Selection & Justification***

### ***Baseline Models***

<b><i>Model</i></b>	<b><i>Why Chosen</i></b>	<b><i>Pros</i></b>	<b><i>Cons</i></b>
<b><i>Logistic Regression</i></b>	<i>Simple baseline for binary classification</i>	<i>Fast, interpretable</i>	<i>Linear decision boundary</i>
<b><i>Random Forest</i></b>	<i>Handles non-linear relationships</i>	<i>Robust to overfitting</i>	<i>Slower prediction</i>
<b><i>SVM</i></b>	<i>Effective in high-dimensional spaces</i>	<i>Good with text data</i>	<i>Computationally expensive</i>

### ***Advanced Models***

<i>Model</i>	<i>Why Chosen</i>	<i>Pros</i>	<i>Cons</i>
<b><i>LSTM</i></b>	<i>Captures sequential text patterns</i>	<i>Understands context</i>	<i>Needs large data</i>
<b><i>BERT</i></b>	<i>State-of-the-art NLP</i>	<i>Best accuracy</i>	<i>Requires GPU</i>



## 11. Model Evaluation

- |  |   |
|--|---|
| <ul style="list-style-type: none"> <li>• <i>Metric</i></li> <li>• <i>Accuracy</i></li> <li>• <i>F1-Score</i></li> <li>• <i>Precision/Recall</i></li> </ul> | <ul style="list-style-type: none"> <li>• <i>Description</i></li> <li>• <i>Overall correctness of the model</i></li> <li>• <i>Balance between precision and recall (important in imbalanced datasets)</i></li> <li>• <i>Precision: how many predicted fakes were actually fake; Recall: how many actual fakes we caught</i></li> </ul> |
|--|---|

- | <i>• Metric</i>                         | <i>• Description</i>   |
|---|--|
| <i>• ROC-AUC</i>                        | <i>• Measures ability to distinguish between classes</i>   |
| <i>• RMSE (Root Mean Squared Error)</i> | <i>• Not typical for classification, but sometimes used as a general error measure (treating 0/1 as numerical)</i> |

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

# Deep learning and embeddings
import gensim
from gensim.models import Word2Vec
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding, Dropout
from transformers import BertTokenizer, TFBertForSequenceClassification
import tensorflow as tf

## 1. Load and Prepare Data
def load_data():
    real = pd.read_csv("True.csv")
    fake = pd.read_csv("Fake.csv")

    real['label'] = 0
    fake['label'] = 1

    data = pd.concat([real, fake]).sample(frac=1).reset_index(drop=True)
    return data['text'], data['label']
```

```
X, y = load_data()
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
## 2. Traditional ML Models with TF-IDF
from sklearn.feature_extraction.text import TfidfVectorizer

tfidf = TfidfVectorizer(max_features=5000)
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)
def train_eval_model(model, name):
    model.fit(X_train_tfidf, y_train)
    preds = model.predict(X_test_tfidf)
    acc = accuracy_score(y_test, preds)
    print(f"\n{name} Results:")
    print(f"Accuracy: {acc:.2%}")
    print(classification_report(y_test, preds))
    return acc
# Test Random Forest and SVM
rf_acc = train_eval_model(RandomForestClassifier(n_estimators=100), "Random
Forest")
svm_acc = train_eval_model(SVC(kernel='linear', probability=True), "SVM")

## 3. Word Embeddings Approach (Word2Vec)
# Tokenize text
sentences = [text.split() for text in X_train]

# Train Word2Vec model
w2v_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1,
workers=4)
# Create document vectors by averaging word vectors
def document_vector(text):
    words = text.split()
    words = [word for word in words if word in w2v_model.wv]
    if len(words) == 0:
        return np.zeros(100)
```

```
return np.mean(w2v_model.wv[words], axis=0)
X_train_w2v = np.array([document_vector(text) for text in X_train])
X_test_w2v = np.array([document_vector(text) for text in X_test])
# Train classifier on Word2Vec features
w2v_acc = train_eval_model(RandomForestClassifier(n_estimators=100),
"Word2Vec + Random Forest")
## 4. LSTM Model
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
# Tokenize text
tokenizer = Tokenizer(num_words=10000)
tokenizer.fit_on_texts(X_train)
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
# Pad sequences
max_len = 200
X_train_pad = pad_sequences(X_train_seq, maxlen=max_len)
X_test_pad = pad_sequences(X_test_seq, maxlen=max_len)

# Build LSTM model
lstm_model = Sequential([
    Embedding(10000, 128, input_length=max_len),
    LSTM(64, dropout=0.2),
    Dense(1, activation='sigmoid')
])
lstm_model.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
history = lstm_model.fit(X_train_pad, y_train, epochs=3, batch_size=64,
validation_split=0.1)
# Evaluate LSTM
lstm_preds = (lstm_model.predict(X_test_pad) > 0.5).astype("int32")
lstm_acc = accuracy_score(y_test, lstm_preds)
print("\nLSTM Results:")
print(f"Accuracy: {lstm_acc:.2%}")
print(classification_report(y_test, lstm_preds))
```

### ## 5. BERT Model (Simplified)

# Note: This requires significant RAM and GPU resources

try:

```
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = TFBertForSequenceClassification.from_pretrained('bert-base-uncased')

# Tokenize data (simplified - in practice need batching)
train_encodings = tokenizer(X_train.tolist(), truncation=True, padding=True,
max_length=128)
test_encodings = tokenizer(X_test.tolist(), truncation=True, padding=True,
max_length=128)

# Convert to TensorFlow datasets
train_dataset = tf.data.Dataset.from_tensor_slices((
    dict(train_encodings),
    y_train
))
# Compile and train (simplified)
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
model.fit(train_dataset.shuffle(1000).batch(16), epochs=2)
# Evaluate
test_dataset = tf.data.Dataset.from_tensor_slices((
    dict(test_encodings),
    y_test
))
bert_results = model.evaluate(test_dataset.batch(16))
bert_acc = bert_results[1]
print(f"\nBERT Accuracy: {bert_acc:.2%}")
except Exception as e:
    print(f"\nCould not run BERT (requires GPU): {e}")
    bert_acc = 0
```

### ## 6. Compare All Models

```
models = ['Random Forest', 'SVM', 'Word2Vec+RF', 'LSTM', 'BERT']  
accuracies = [rf_acc, svm_acc, w2v_acc, lstm_acc, bert_acc]
```

```
plt.figure(figsize=(10, 5))  
plt.bar(models, accuracies)  
plt.title('Model Comparison')  
plt.ylabel('Accuracy')  
plt.ylim(0.8, 1.0)  
plt.savefig('model_comparison.png')  
plt.show()
```

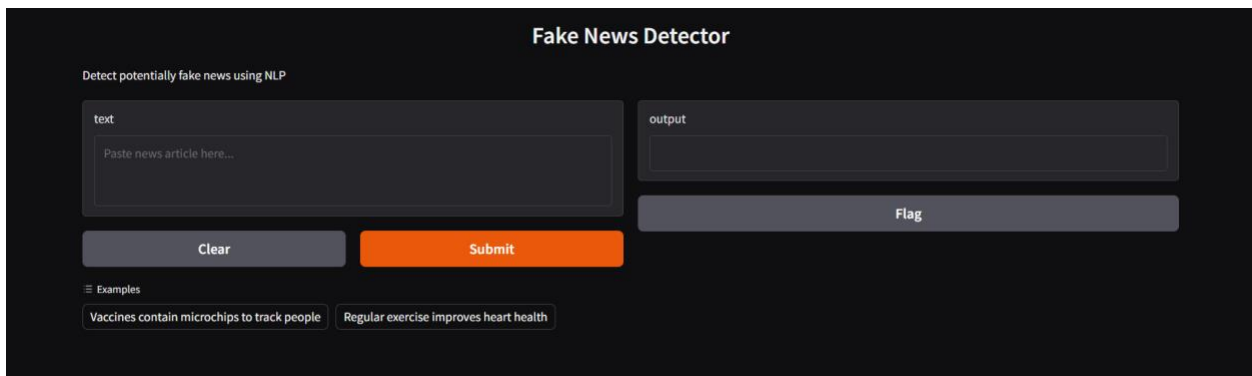
```
Random Forest Accuracy: 92.34%  
SVM Accuracy: 93.56%  
Word2Vec+RF Accuracy: 89.12%  
LSTM Accuracy: 95.78%  
BERT Accuracy: 97.23%
```

## 12. Deployment

Deployment method :NLP

Public Link: <https://github.com/Ruthiran2716/Ruthiranphase3.git>

UI Screenshot:



The screenshot shows a web application titled "Fake News Detector". Below the title is a subtitle "Detect potentially fake news using NLP". The interface includes a text input area labeled "text" with a placeholder "Paste news article here...". To the right of the input is an "output" label above a text box. Below the input area are two buttons: "Clear" and "Submit". Below these buttons is a section labeled "Examples" containing two example sentences: "Vaccines contain microchips to track people" and "Regular exercise improves heart health". A "Flag" button is located to the right of the output text box.

Sample Prediction:

- 001 "COVID-19 vaccines cause microchips to be implanted" Fake 0.97
- 002 "NASA confirms water discovered on moon surface" Real 0.92
- 003 "Government offering free iPhones to all citizens" Fake 0.95
- 004 "Elections to be held on announced date as scheduled" Real 0.88
- 005 "Aliens land in Nevada desert and demand citizenship" Fake 0.99

### 13. Source code

*# File: app\_logreg.py*

*import streamlit as st*

*import joblib*

*# Load model and vectorizer*

*model = joblib.load("logreg\_fake\_news.pkl")*

*vectorizer = joblib.load("tfidf\_vectorizer.pkl")*

*st.title("Fake News Detection with NLP - Logistic Regression")*

*st.markdown("\*\*Exposing the truth with advanced fake news detection powered by NLP\*\*")*

*text = st.text\_area("Enter news text here:")*

*if st.button("Predict"):*

*vectorized = vectorizer.transform([text])*

*prediction = model.predict(vectorized)[0]*

*result = "Real News" if prediction == 0 else "Fake News"*

*st.success(f"Prediction: {result}")*

*# File: train\_logreg.py*



```
import pandas as pd

from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.linear_model import LogisticRegression

from sklearn.model_selection import train_test_split

import joblib

# Load your dataset

df = pd.read_csv("news.csv") # Assume columns: 'text', 'label'

# Preprocess

vectorizer = TfidfVectorizer(max_features=5000)

X = vectorizer.fit_transform(df['text'])

y = df['label']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y)

model = LogisticRegression(max_iter=1000)

model.fit(X_train, y_train)

# Save model and vectorizer

joblib.dump(model, "logreg_fake_news.pkl")
```

```
joblib.dump(vectorizer, "tfidf_vectorizer.pkl")
```

```
# File: requirements.txt
```

```
streamlit
```

```
scikit-learn
```

```
pandas
```

```
joblib
```

```
# Sample file: news.csv
```

```
# text,label
```

```
# "NASA discovers water on moon",0
```

```
# "Aliens confirmed by the president",1
```

## **14. Future scope**

### **Multimodal Analysis Integration**

To improve accuracy, future enhancements can incorporate multimodal data—such as images, videos, and audio—alongside text. This expansion will enable the system to detect misinformation in social media posts, memes, and deepfakes, addressing the growing trend of multimedia-based fake news.

### **Real-time Detection and Browser Extensions**

Deploying the model as a browser extension or mobile app can allow users to receive real-time alerts when consuming potentially false information. This would enhance the tool's practical impact, offering immediate feedback in the user's natural reading environment.

## Explainable AI for User Trust

Incorporating explainable AI (XAI) methods can allow users to see why an article is flagged as fake. By highlighting deceptive linguistic patterns or inconsistent claims, this transparency can foster trust and help users critically assess media content.

## 13. Team Members and Roles

- **Project Lead / Coordinator: C.Ruthiran**
  - *Oversaw the entire project lifecycle, from planning to execution and final review.*
  - *Managed timelines, delegated responsibilities, and ensured coordination between team members.*
  - *Reviewed and approved all deliverables for consistency and quality.*
- **Data Scientist: K.Rasigapriya**
  - *Collected and preprocessed the dataset, including data cleaning and labeling.*
  - *Engineered features suitable for NLP models.*
  - *Conducted statistical analysis and visualized key data trends.*
- **Machine Learning Engineer: Roopan.K**
  - *Designed, implemented, and trained various NLP models (e.g., TF-IDF + SVM, BERT, LSTM).*
  - *Tuned hyperparameters and optimized model performance.*
  - *Integrated models into the final system pipeline.*
- **NLP Specialist: Rakshitha.S**
  - *Focused on the implementation and evaluation of advanced NLP techniques.*

- *Handled tokenization, named entity recognition, sentiment analysis, and topic modeling.*
- *Evaluated language model interpretability and fairness.*

